

Vision Transformer – A Deep Learning model for Image Analysis and Understanding

CML: Control, Machine Learning and Numerics

Group 29



Apurwa Agrawal

Masters in Data Science
Mtrikel Nr. 23161645



Kshitij Dua

Masters in Data Science
Mtrikel Nr. 23032110



Sneha Kumari

Masters in Data Science
Mtrikel Nr. 22970359



Amit Jadhav

Masters in Data Science
Mtrikel Nr. 23036144



Rajesh Madhipati

Masters in Data Science
Mtrikel Nr. 23007086

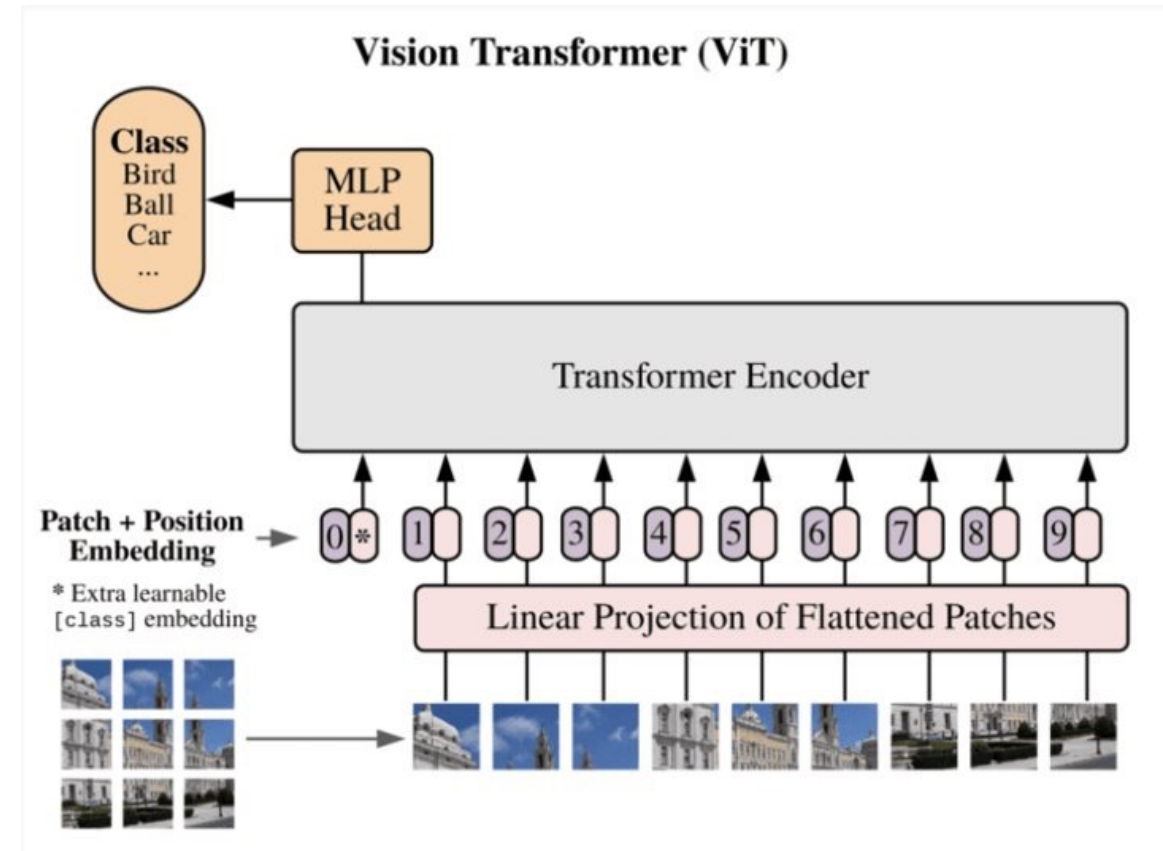
Table of Content

01	Introduction	06	VIT vs CNN
02	Motivation of Vision Transformers	07	Limitations of ViT
03	Evolution of Transformers	08	Discussion and Conclusion
04	Vision Transformer Architecture	09	Future Improvements
05	Training Vision Transformers	10	References

Introduction

Introduction


- Ability to extract meaningful insights from images is crucial for driving advancements across various industries and disciplines.
- CNNs faces problem in **capturing long-range dependencies** and modelling global context within images.
- ViT extends the success of the Transformer architecture.
- ViT understand **complex relationships** within images and make informed predictions.



Motivation

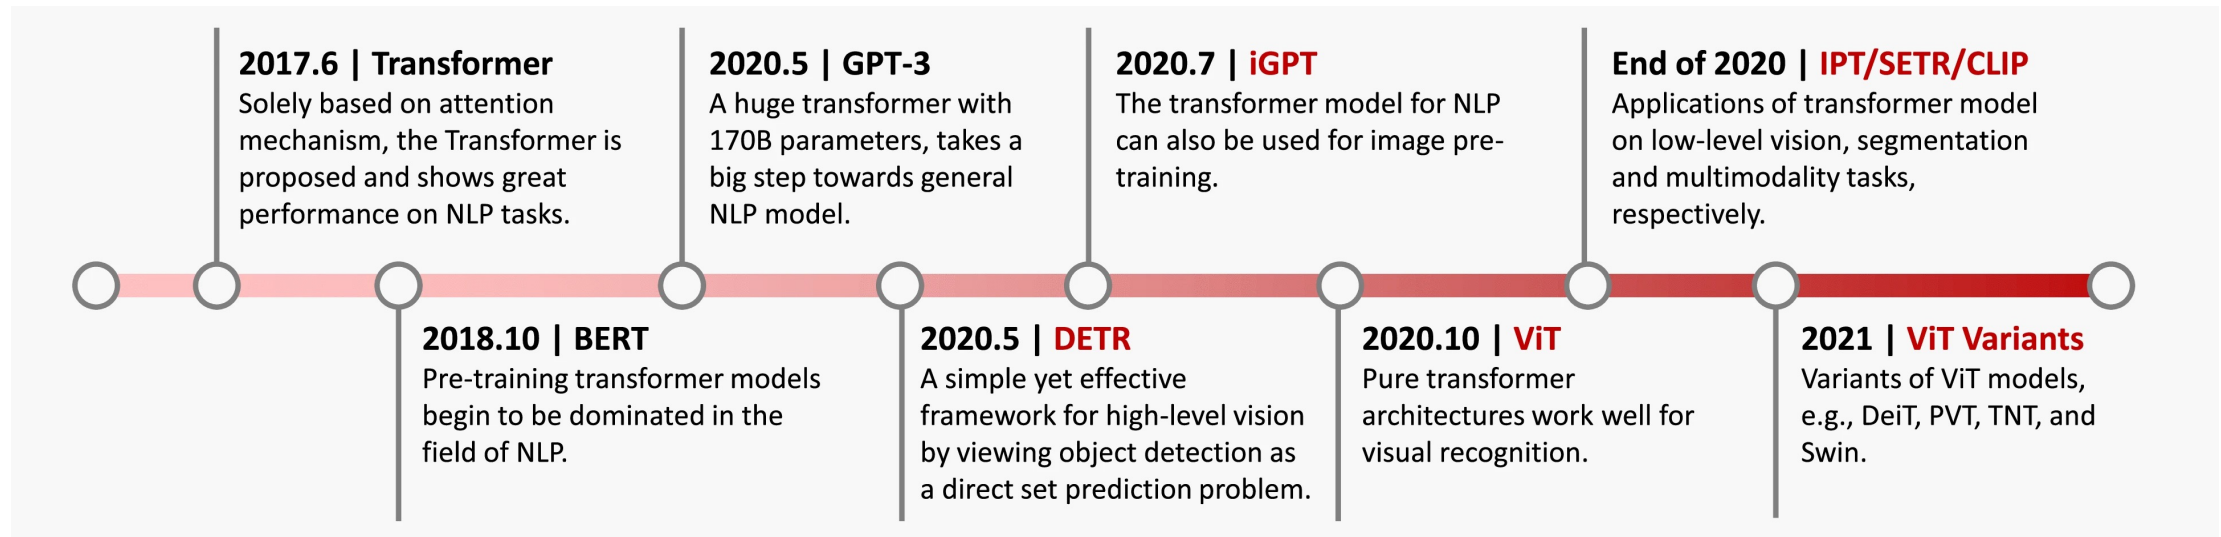
Motivation

- Capturing Long-Range Dependencies.
- Flexibility with Variable-Sized Inputs.
- Transferability and Generalization.
- Learning Abstract Representations.



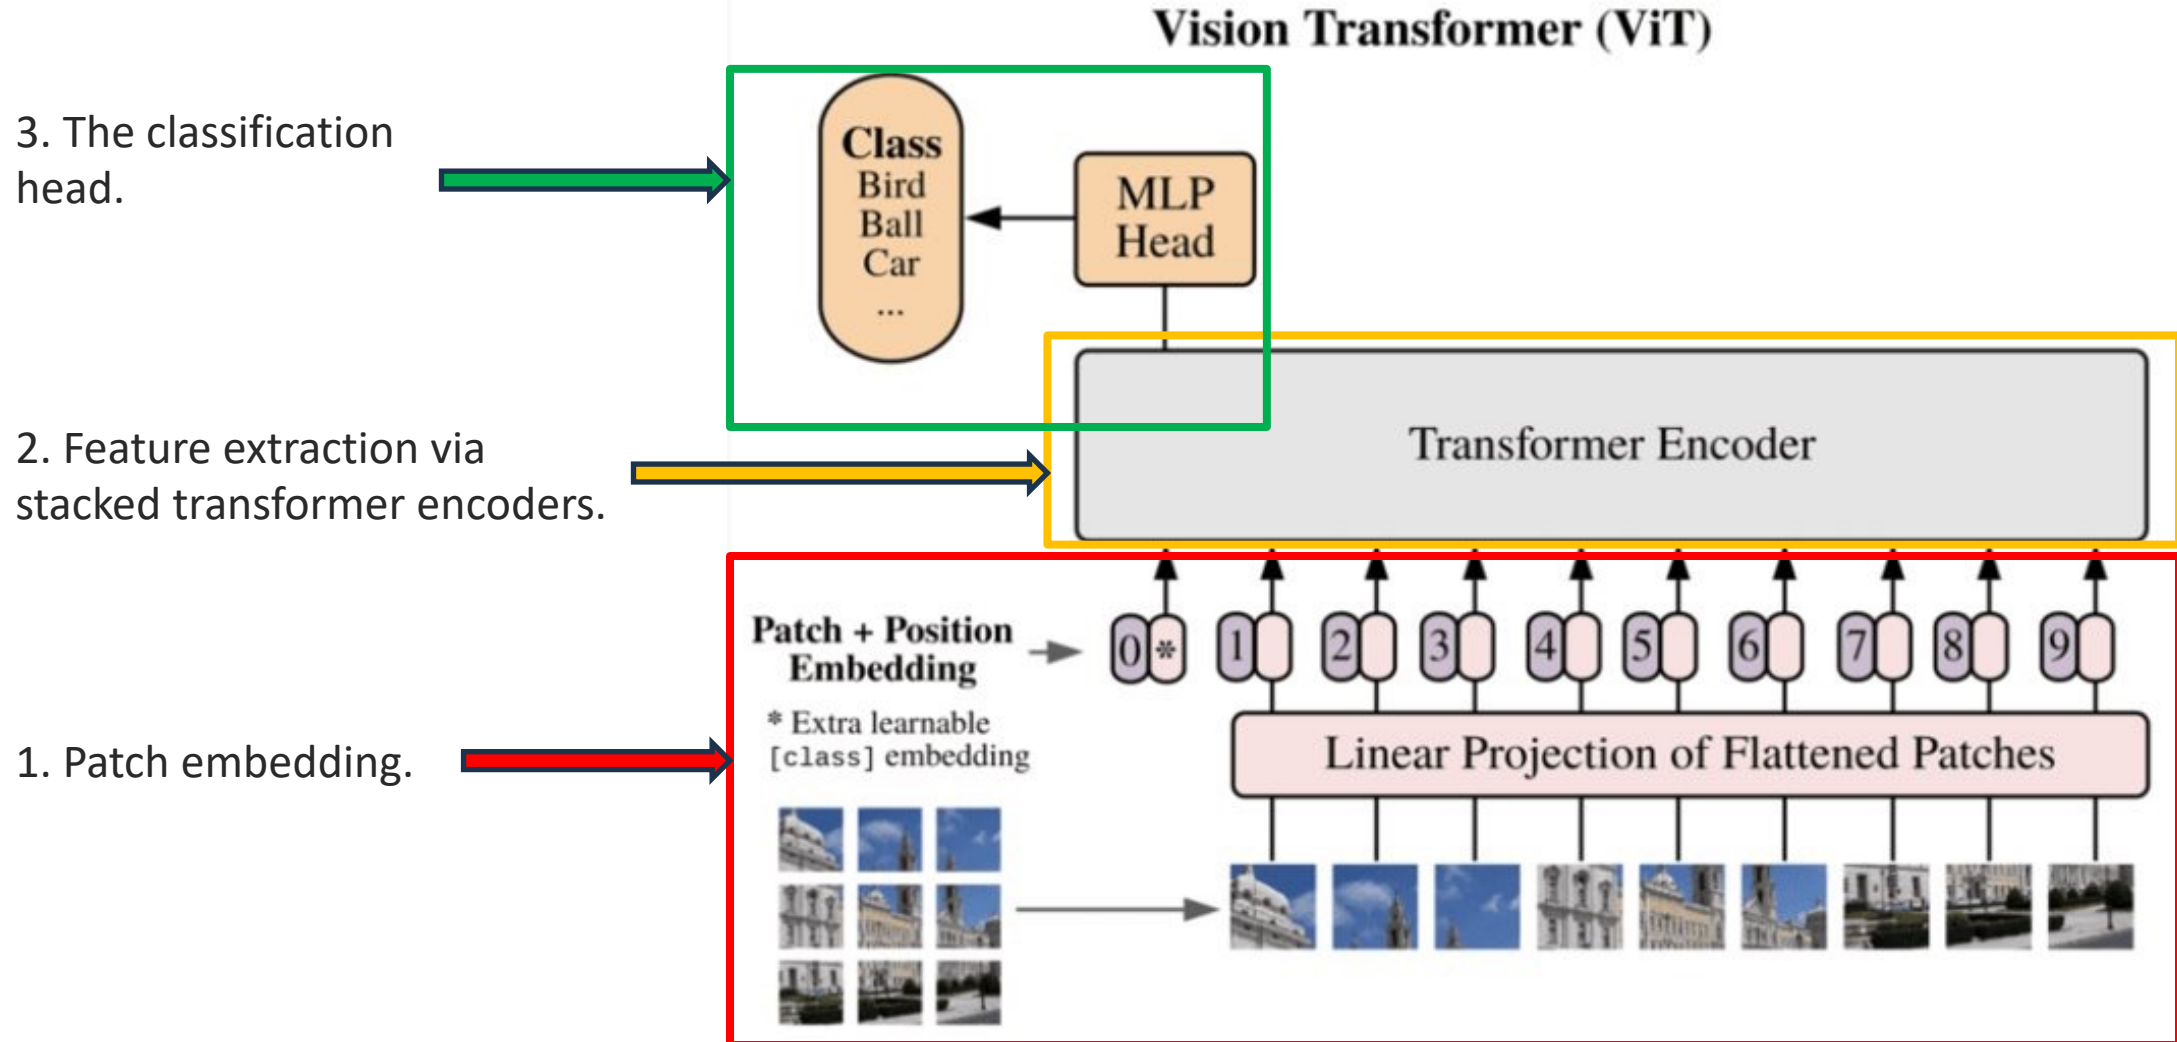
Evolution of Transformers

Evolution of Transformers

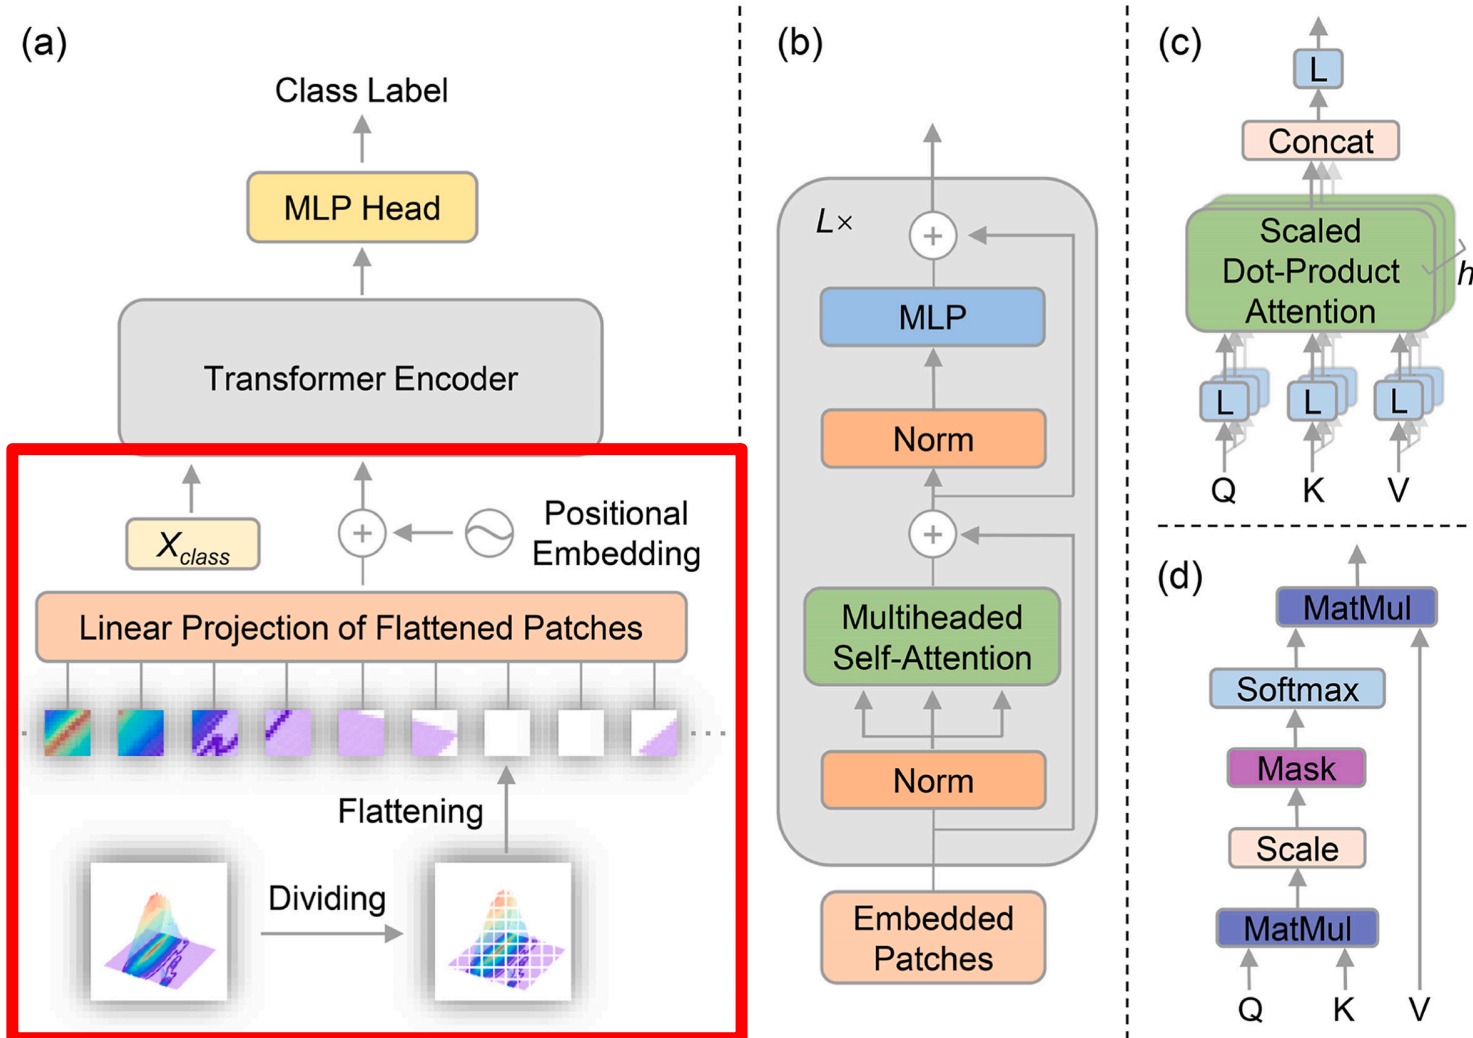


Vision Transformer Architecture

3 Parts of Architecture



1. Patch Embedding



1. Patch Embedding



Batch of
images \mathbf{b}

1. Patch Embedding



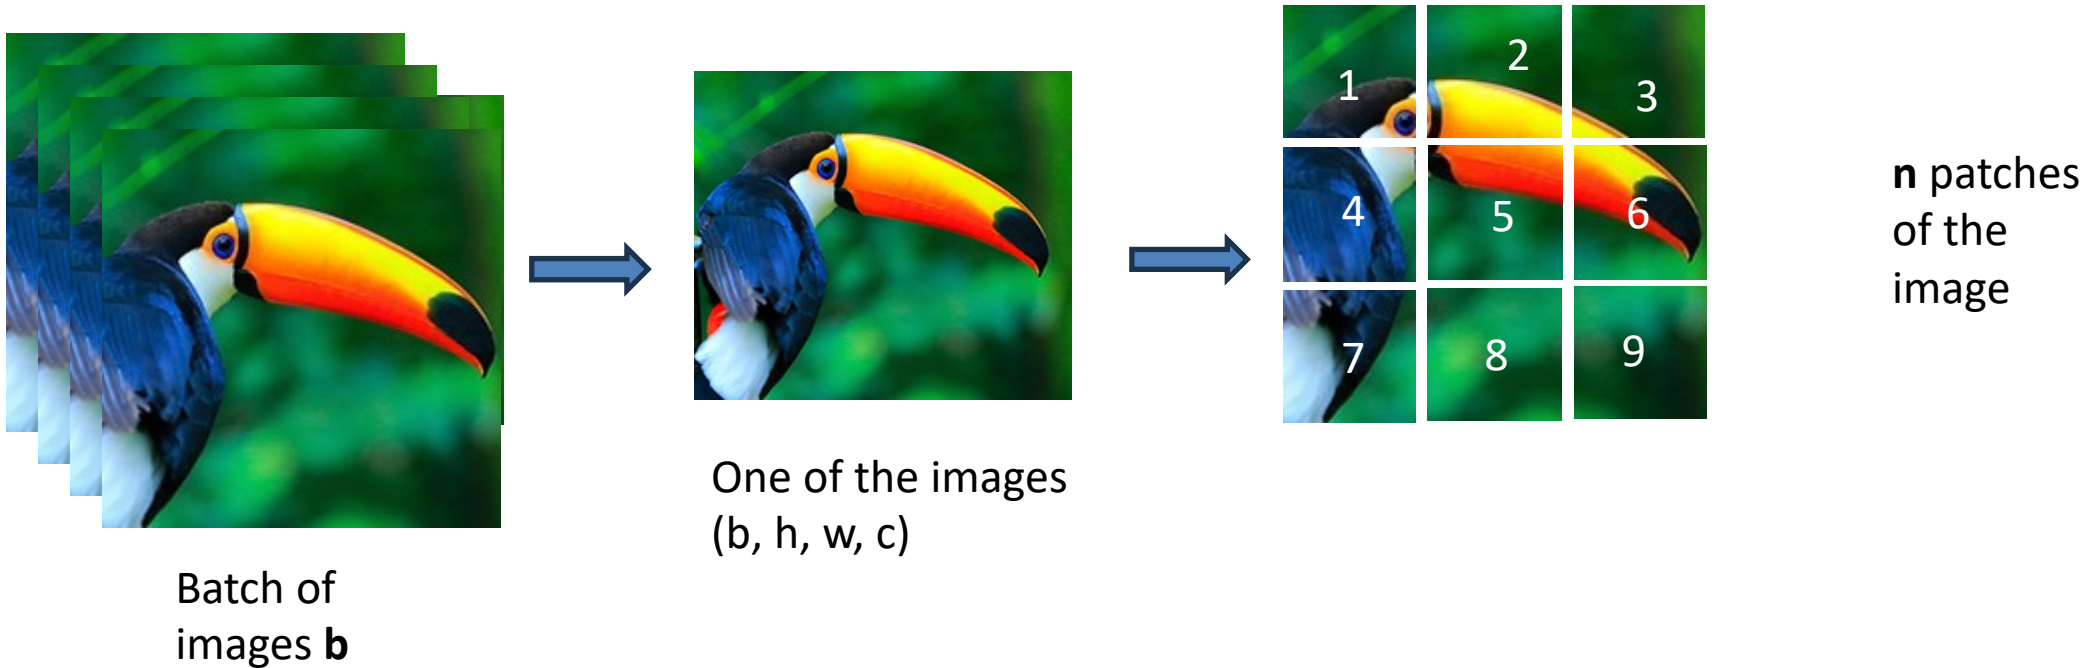
Batch of
images \mathbf{b}



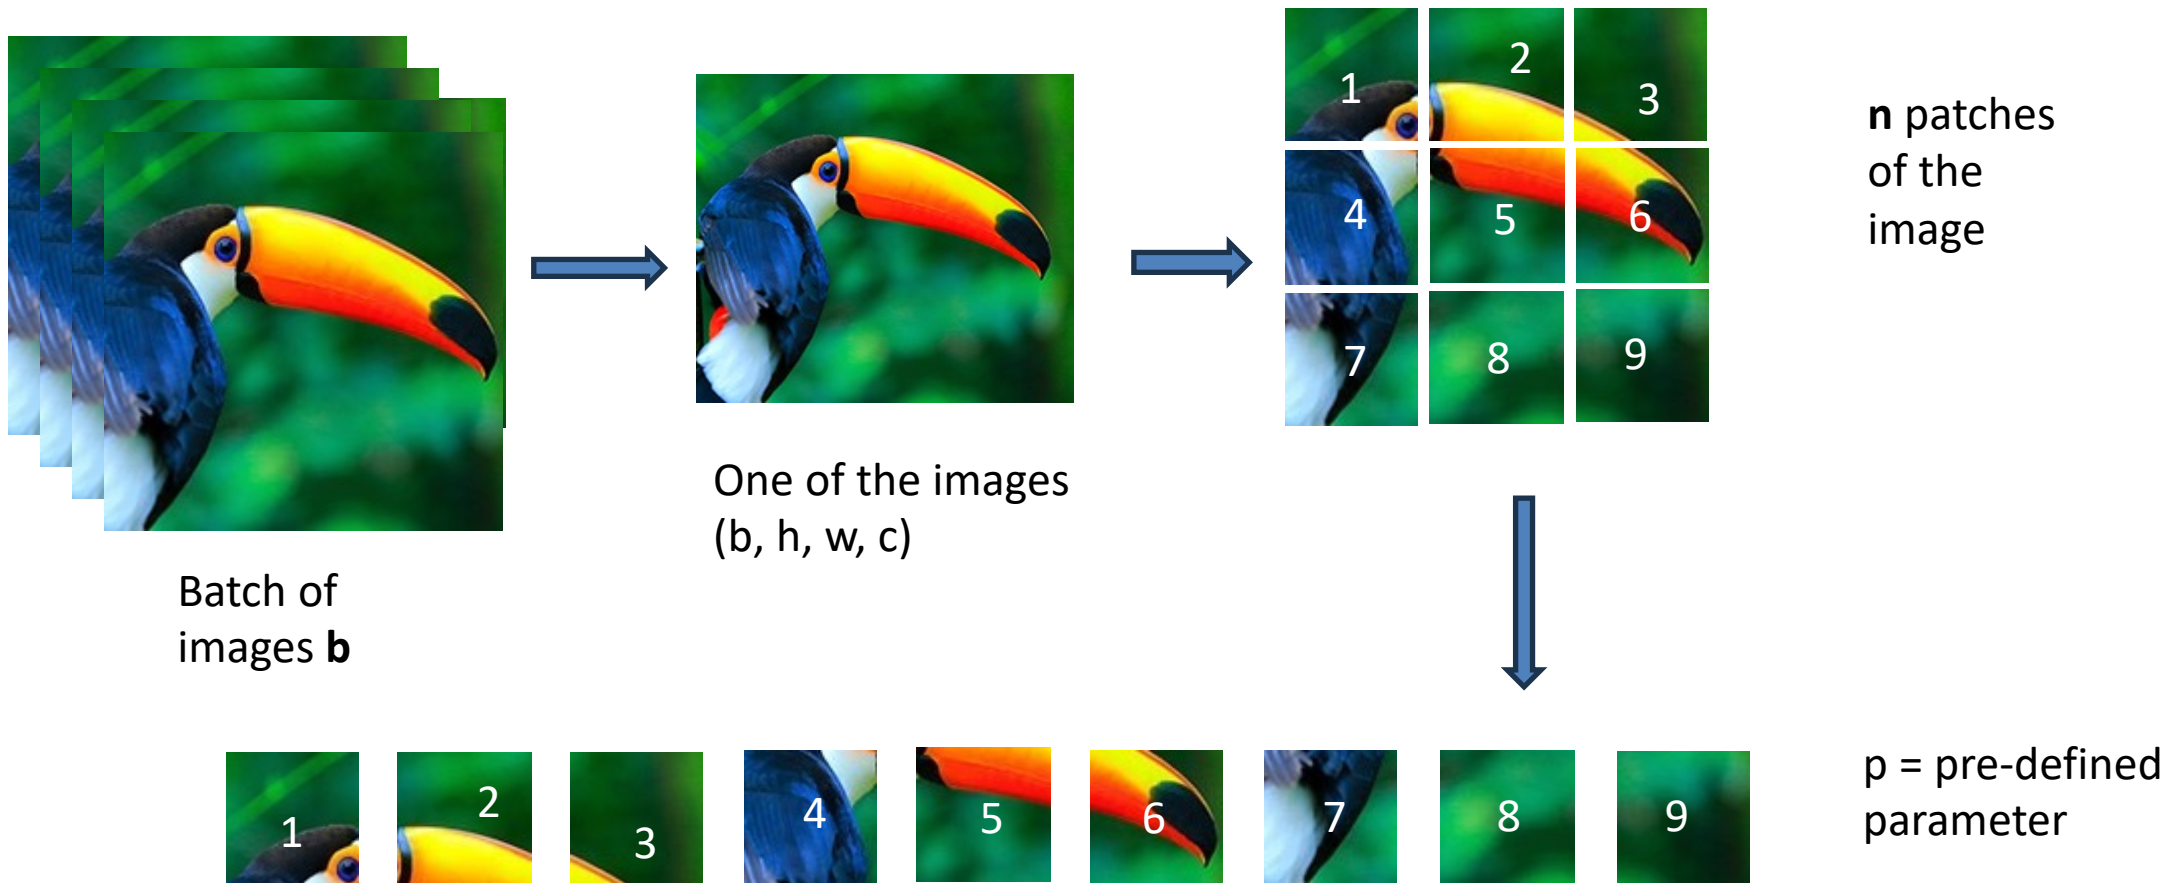
One of the images
(b, h, w, c)

b = batch
 h = height
 w = width
 c = channel

1. Patch Embedding

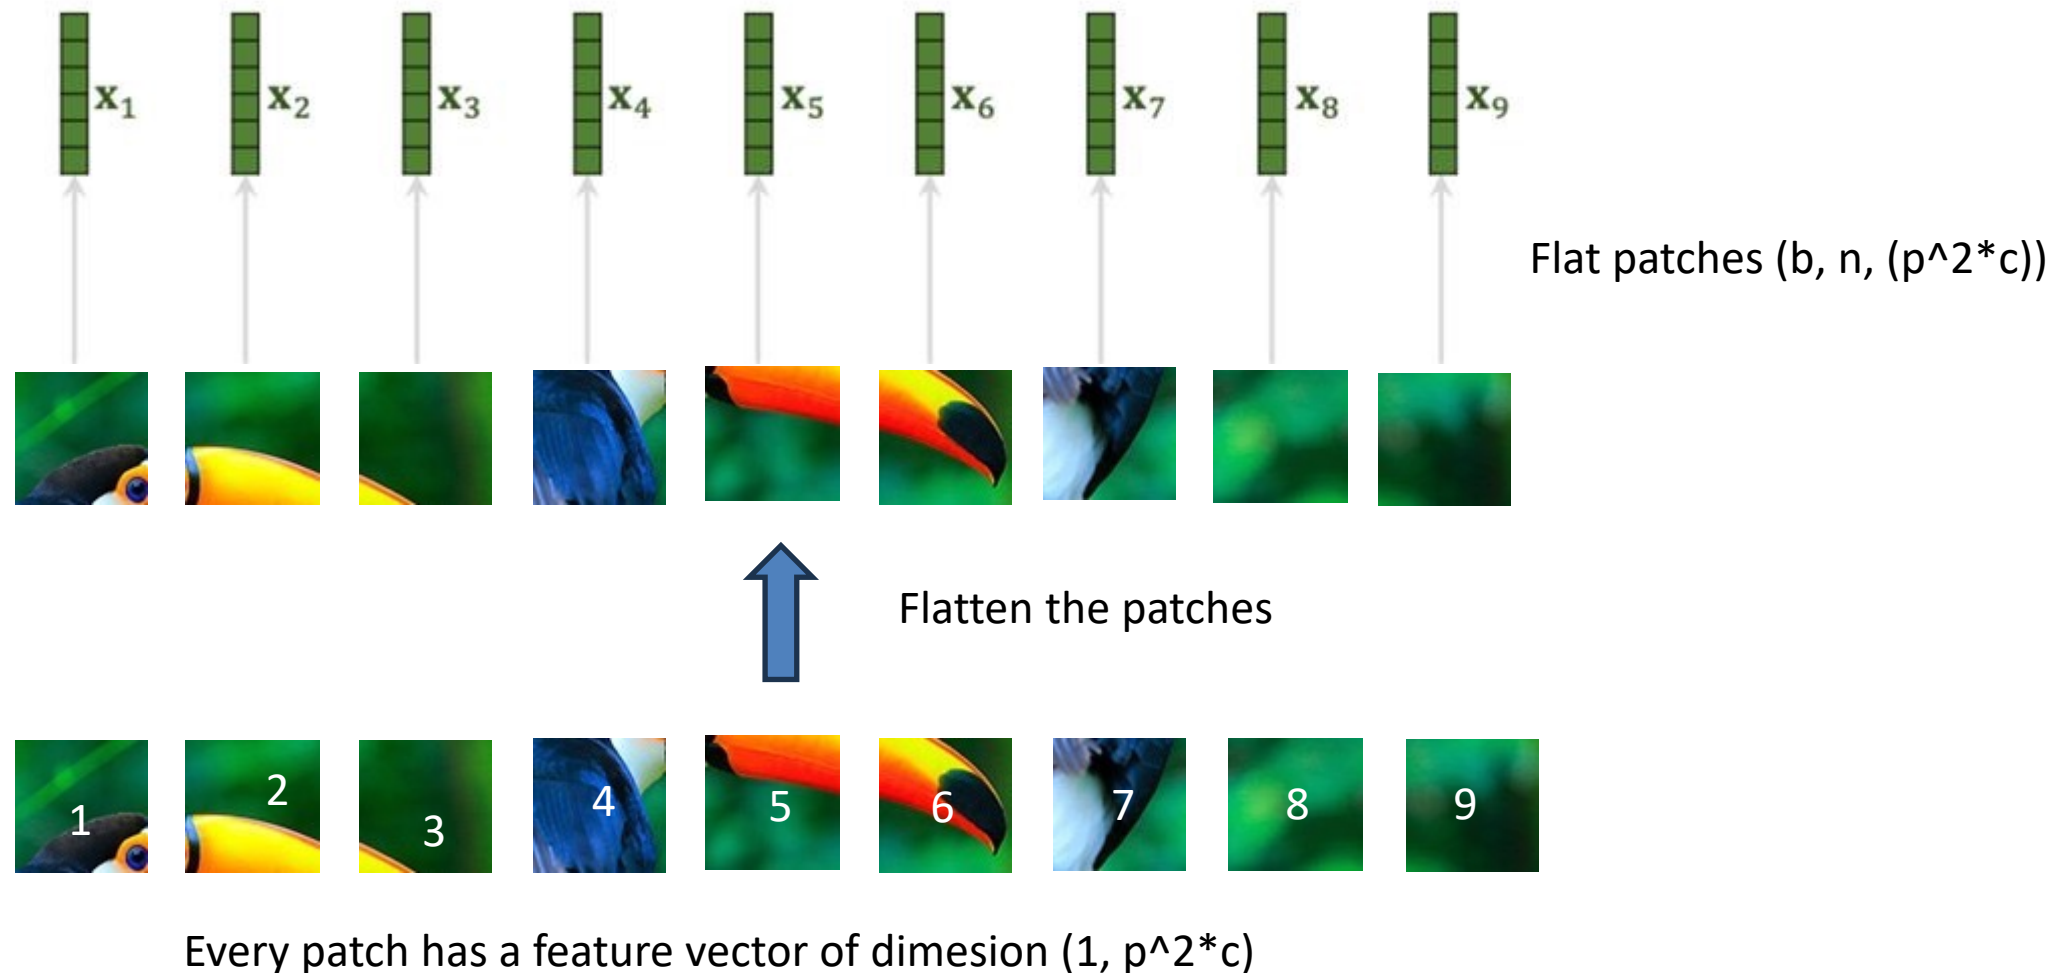


1. Patch Embedding




The image is split into n square patches of shape (p, p, c) ,

1. Patch Embedding



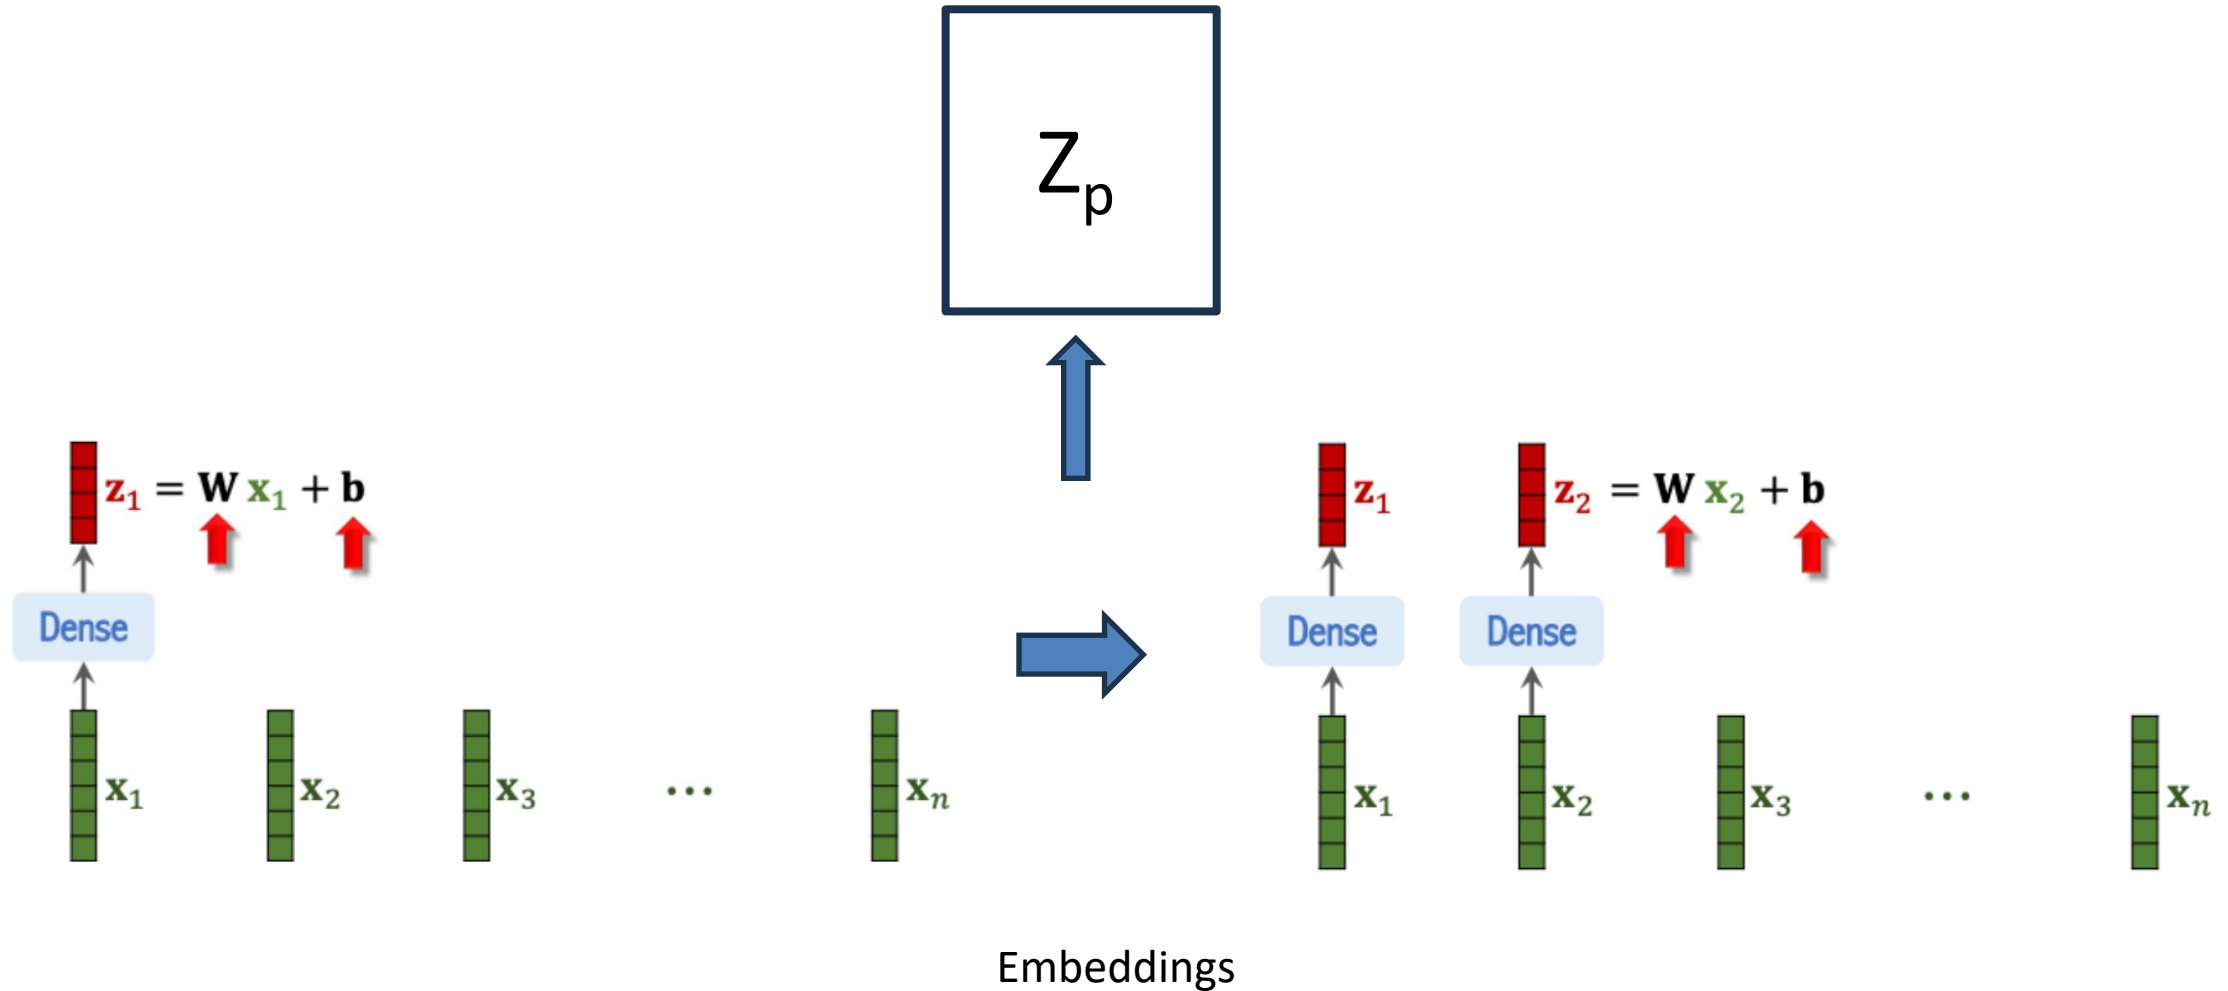
1. Patch Embedding




Embeddings $((p^2 * c), d)$

The flattened patches are multiplied with a **trainable** embedding tensor, which learns to linearly project each flat patch

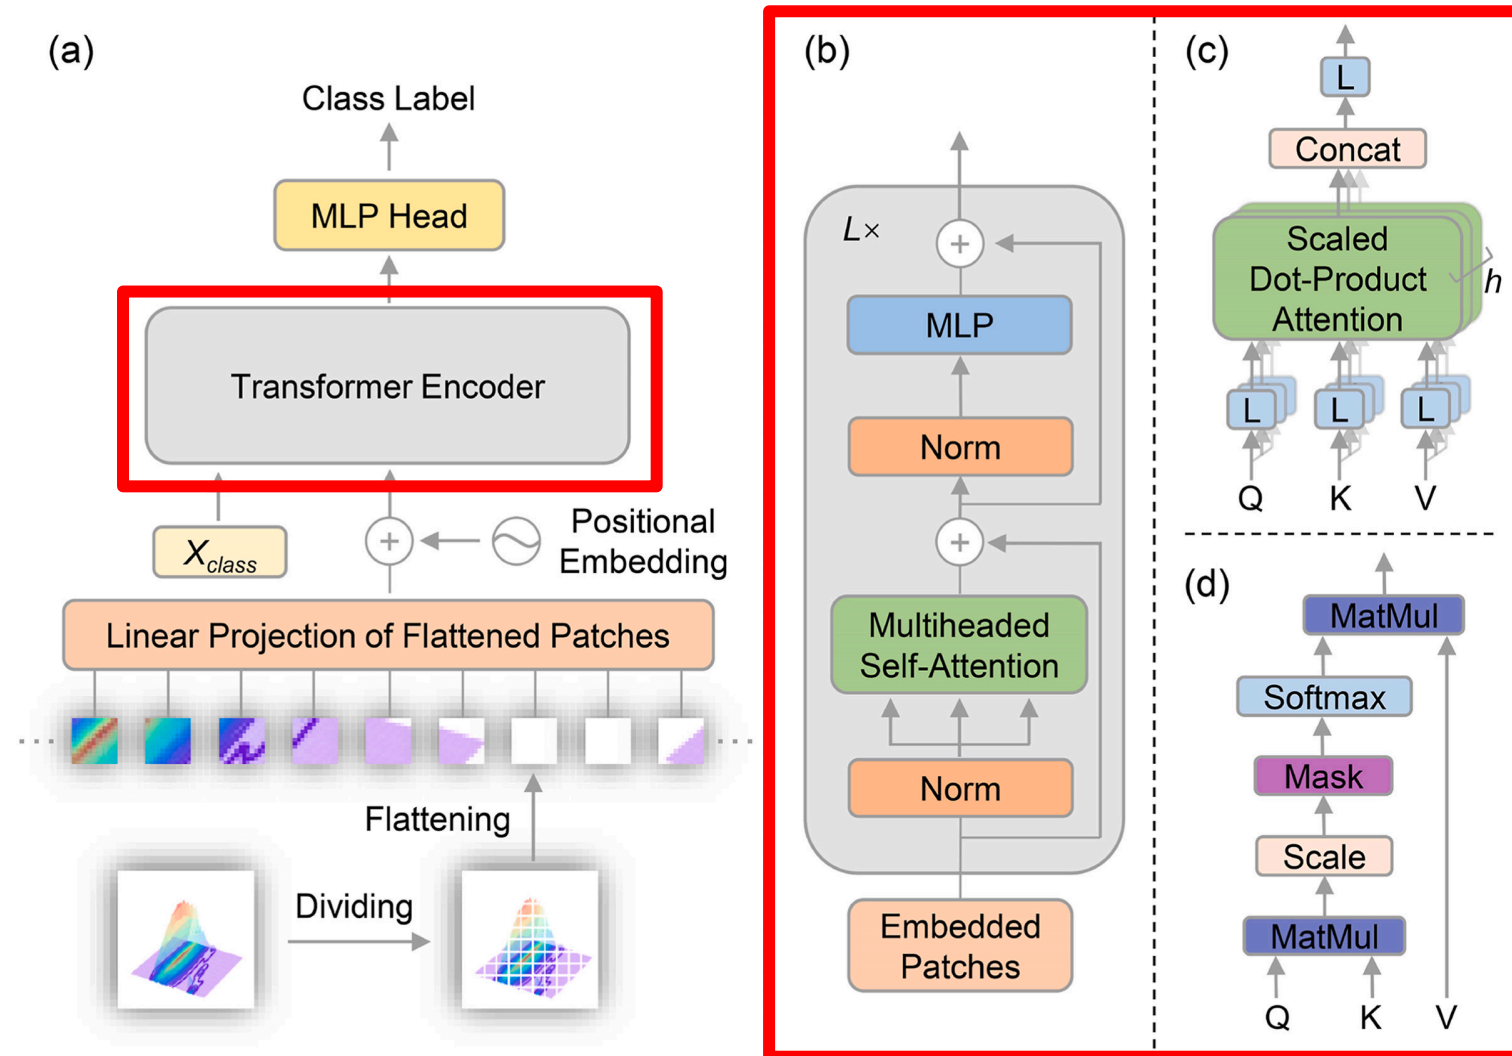
1. Patch Embedding



1. Patch Embedding




2. Transformer Encoding



2.1 Layer Normalization


2. Transformer Encoding

Batch Normalization



Sample for all training examples

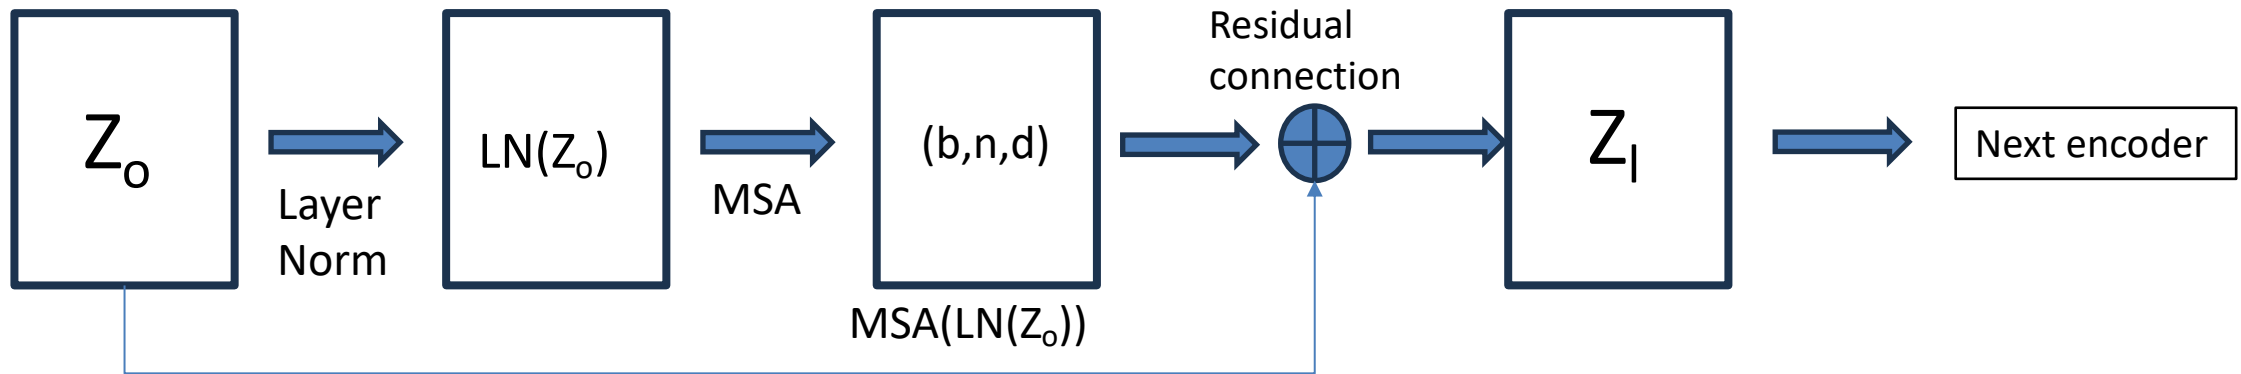
Layer Normalization



Same for all features dimensions

2.1 Layer Normalization

2. Transformer Encoding



Layer Normalization:

$$\mu_i = \frac{1}{n} \sum_{j=1}^n z_{ij}$$

$$\sigma_i^2 = \frac{1}{n} \sum_{j=1}^n (z_{ij} - \mu_i)^2$$

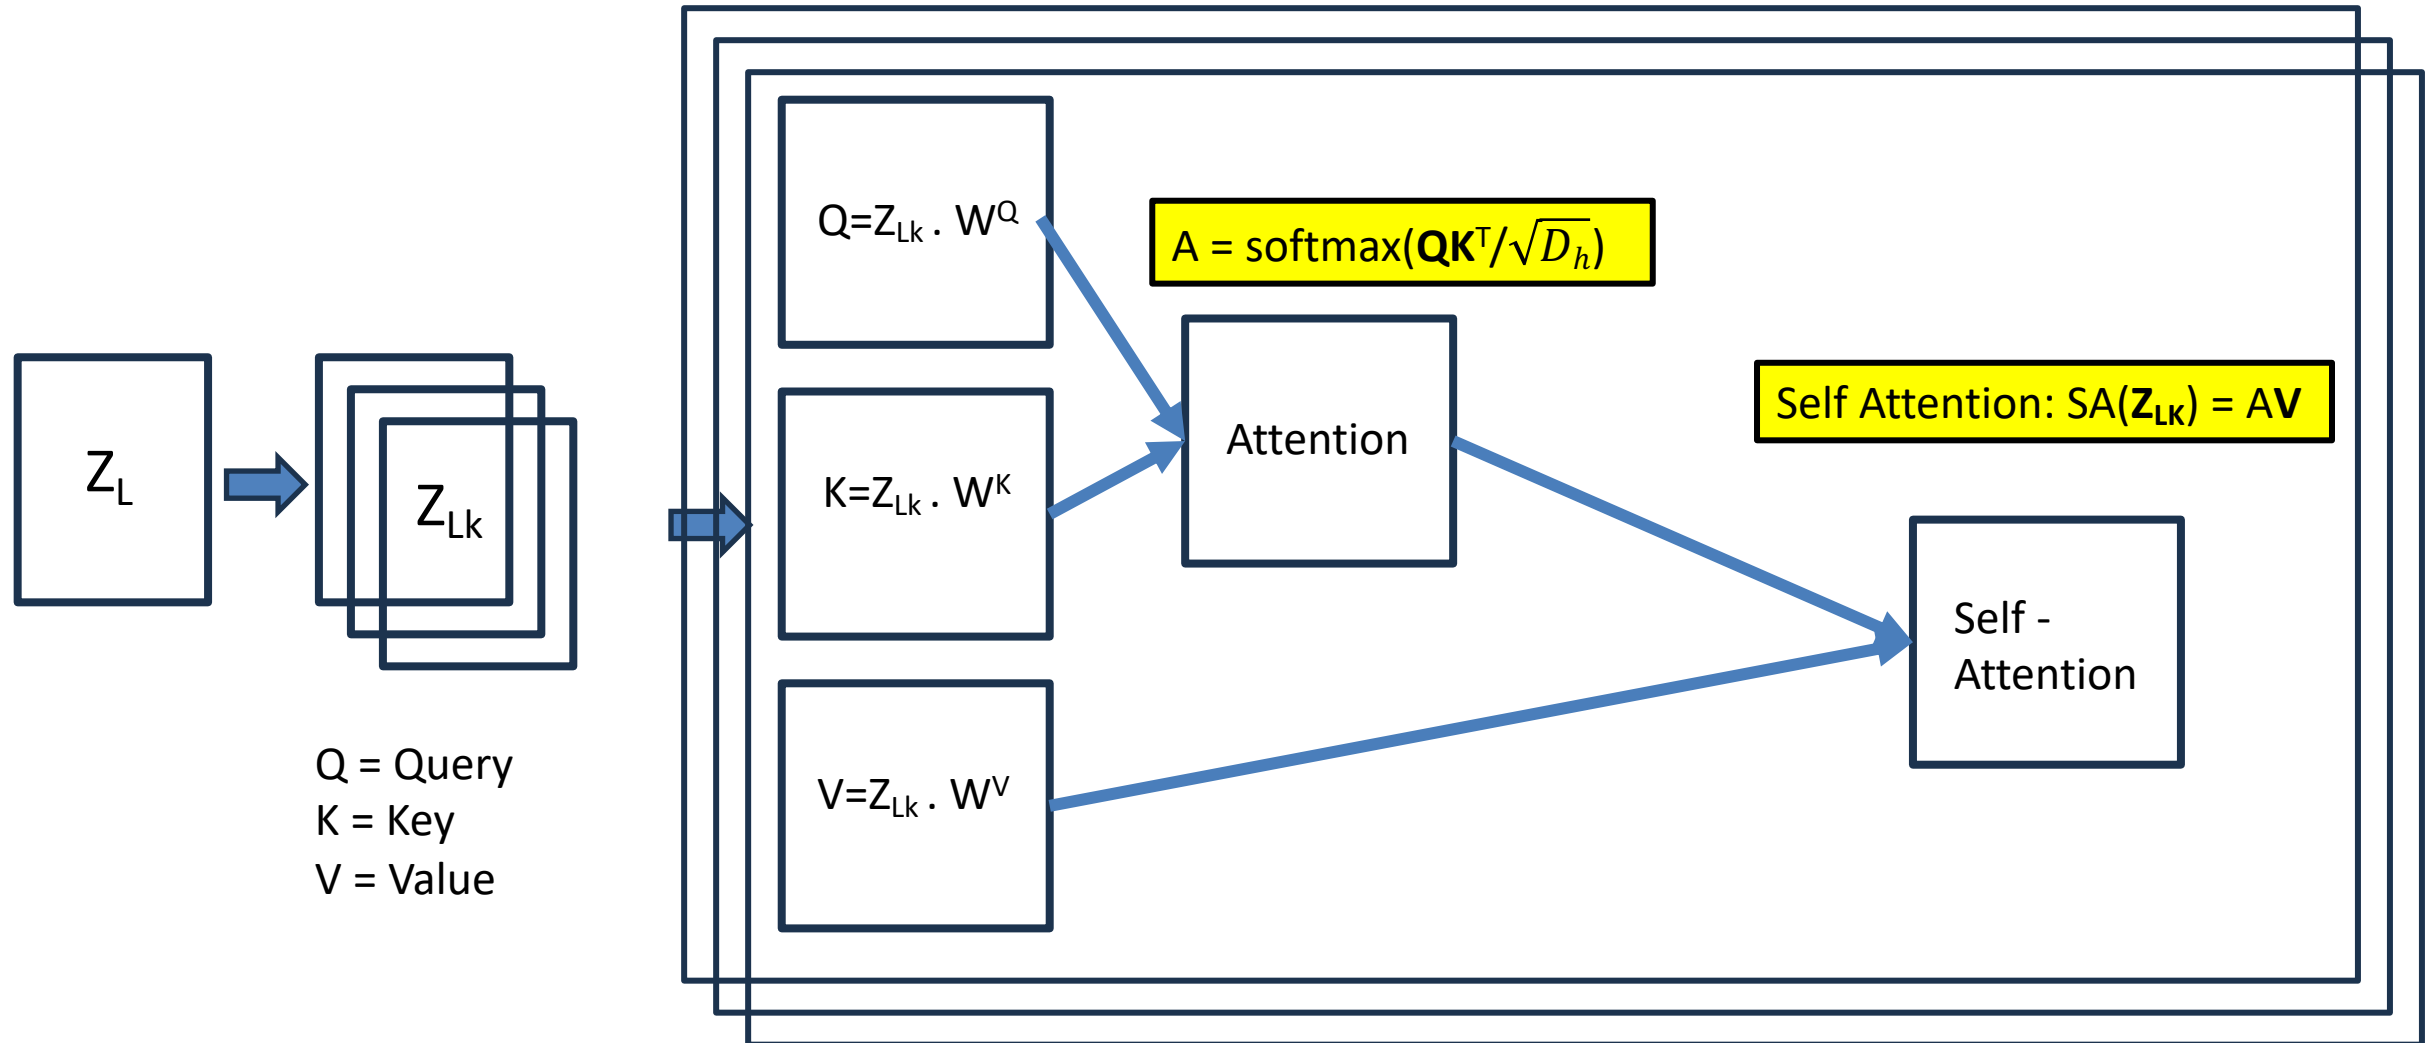
$$\hat{X}_{ij} = \frac{z_{ij} - \mu_i}{\sqrt{\sigma_i^2 + \epsilon}}$$

$$LN(Z_o) = \hat{X}$$

$$Z'_e = MSA(LN(Z_o)) + Z_l, \text{ where } e = 1 \dots L \quad (2)$$

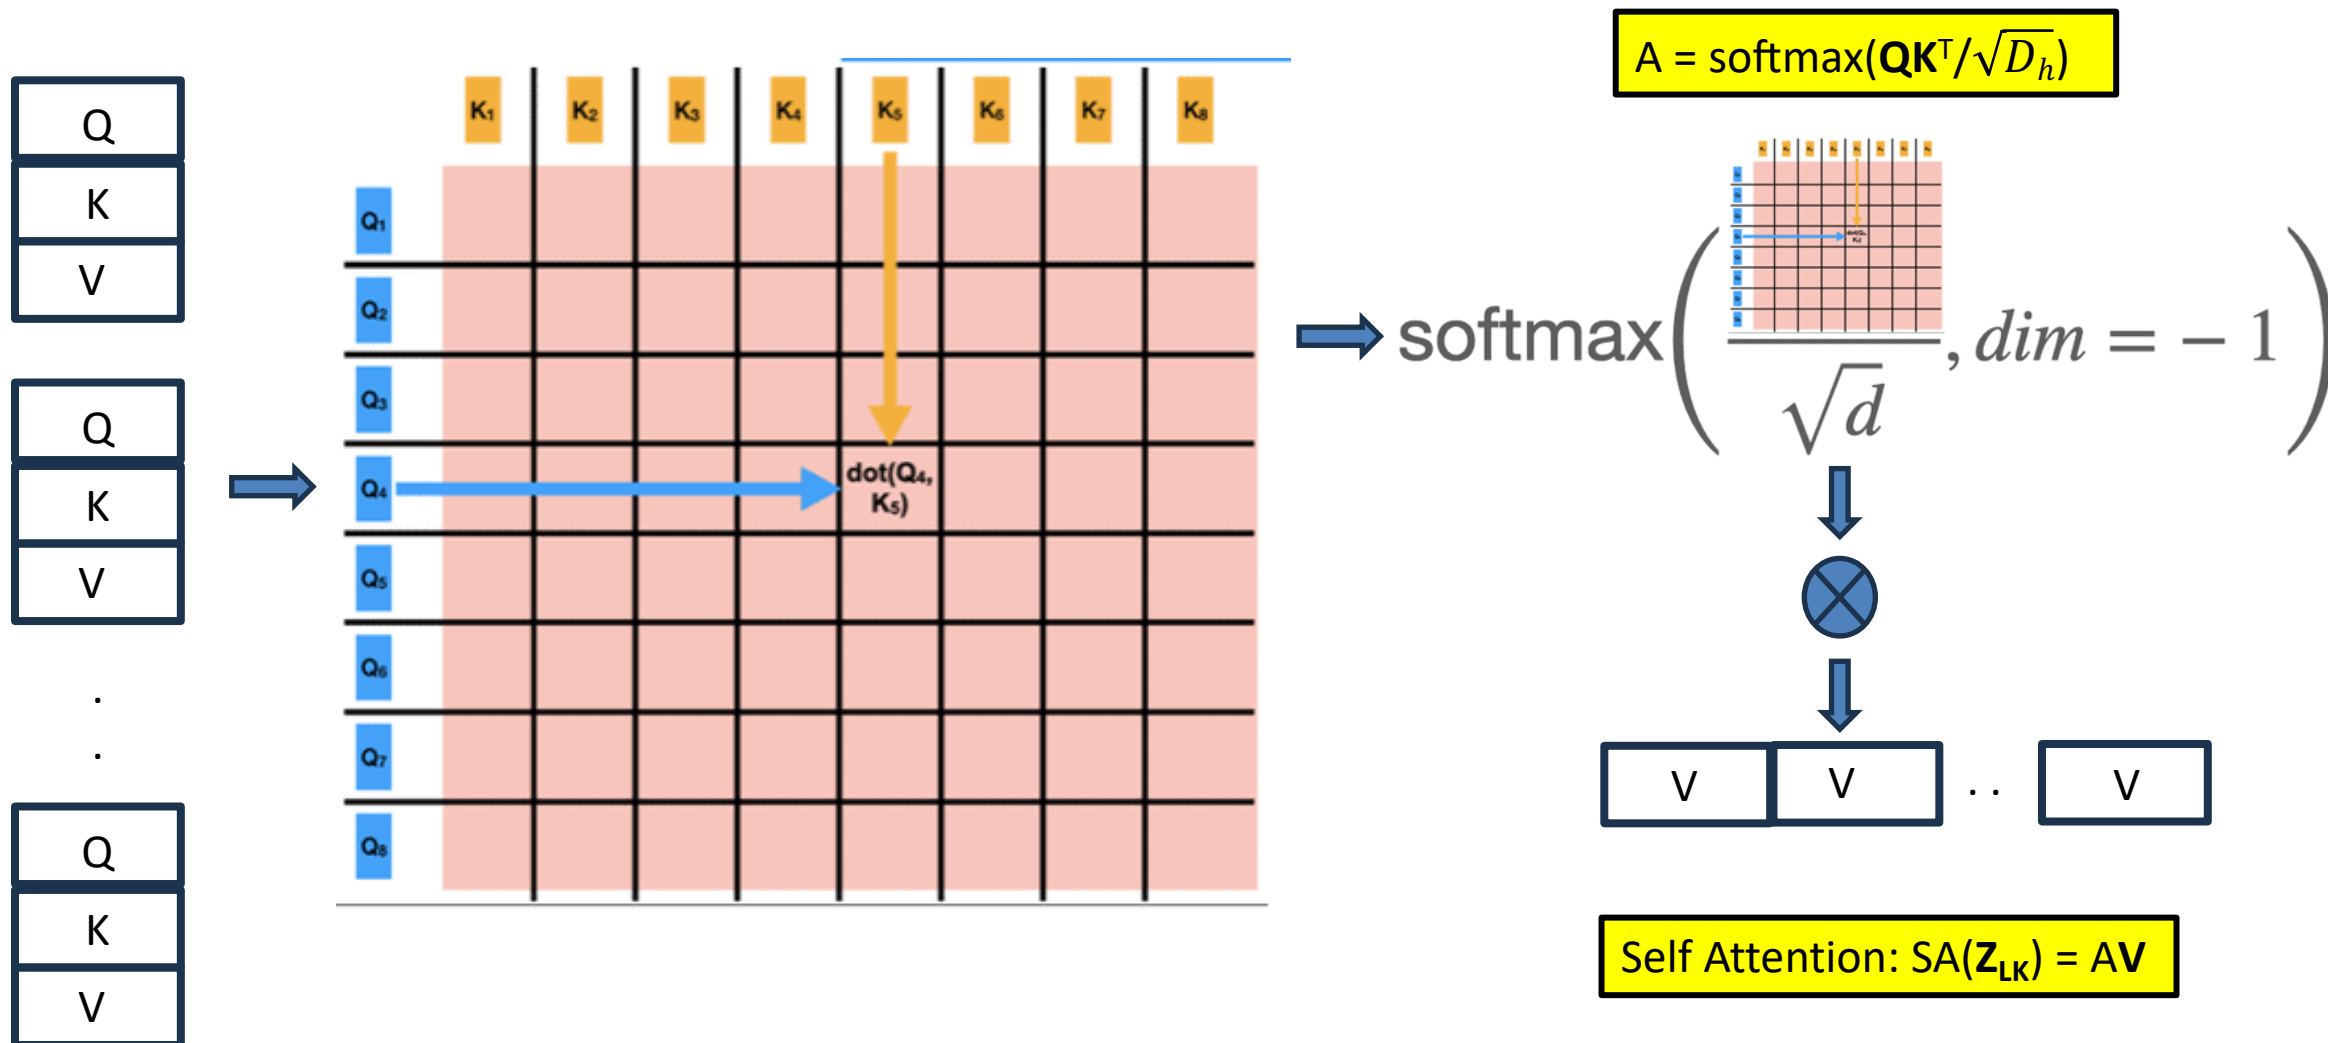
2.2 Multi-head Attention

2. Transformer Encoding



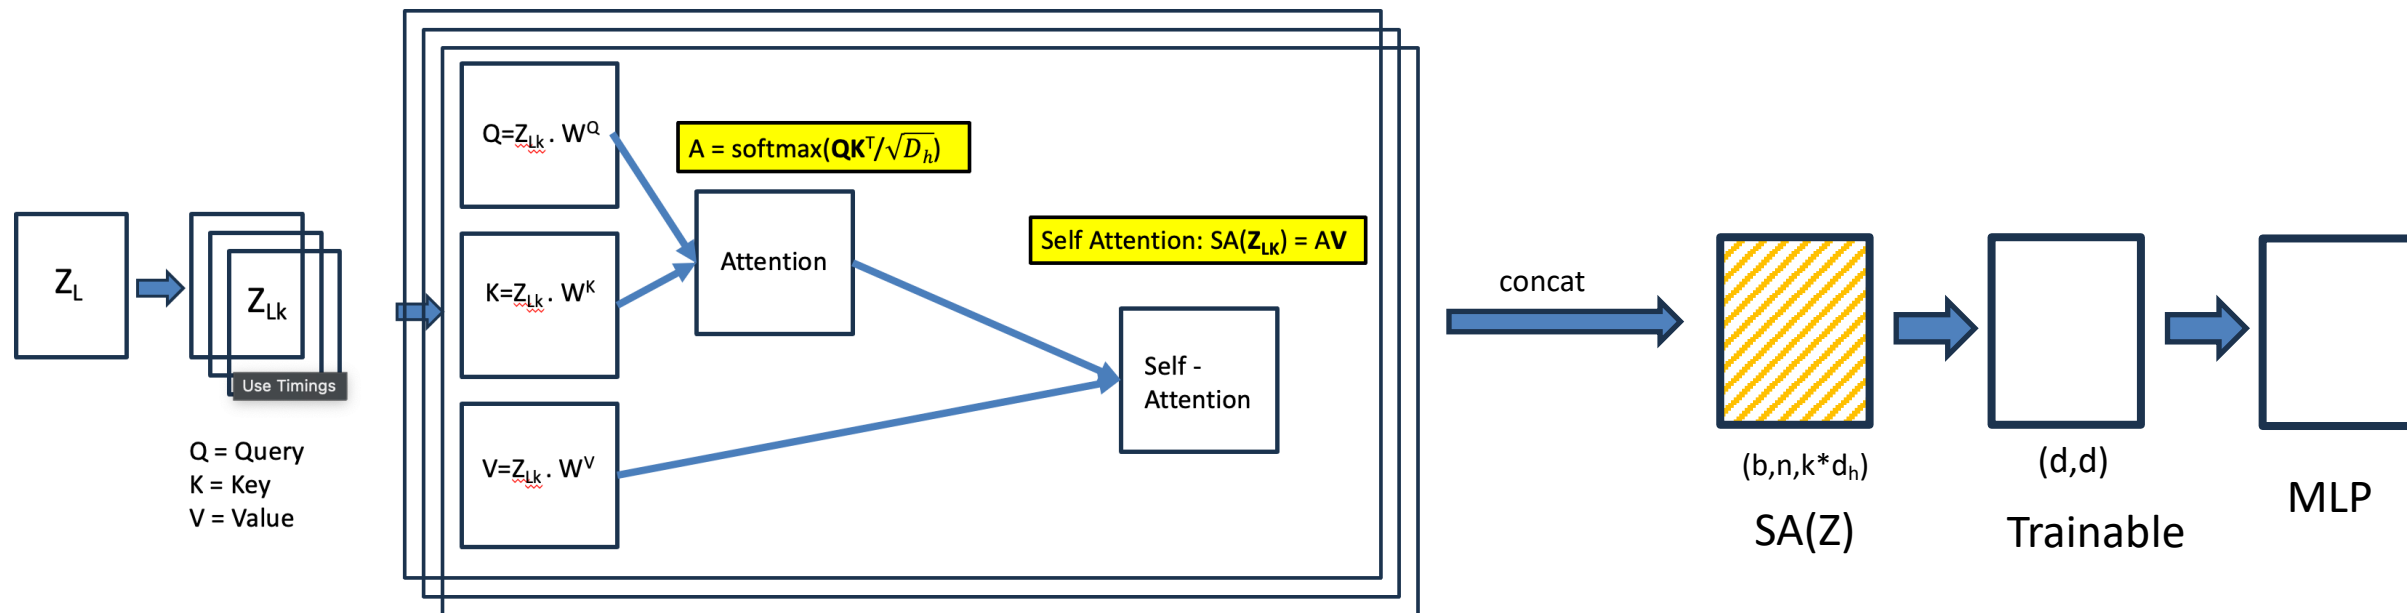
2.2 Multi-head Attention

2. Transformer Encoding

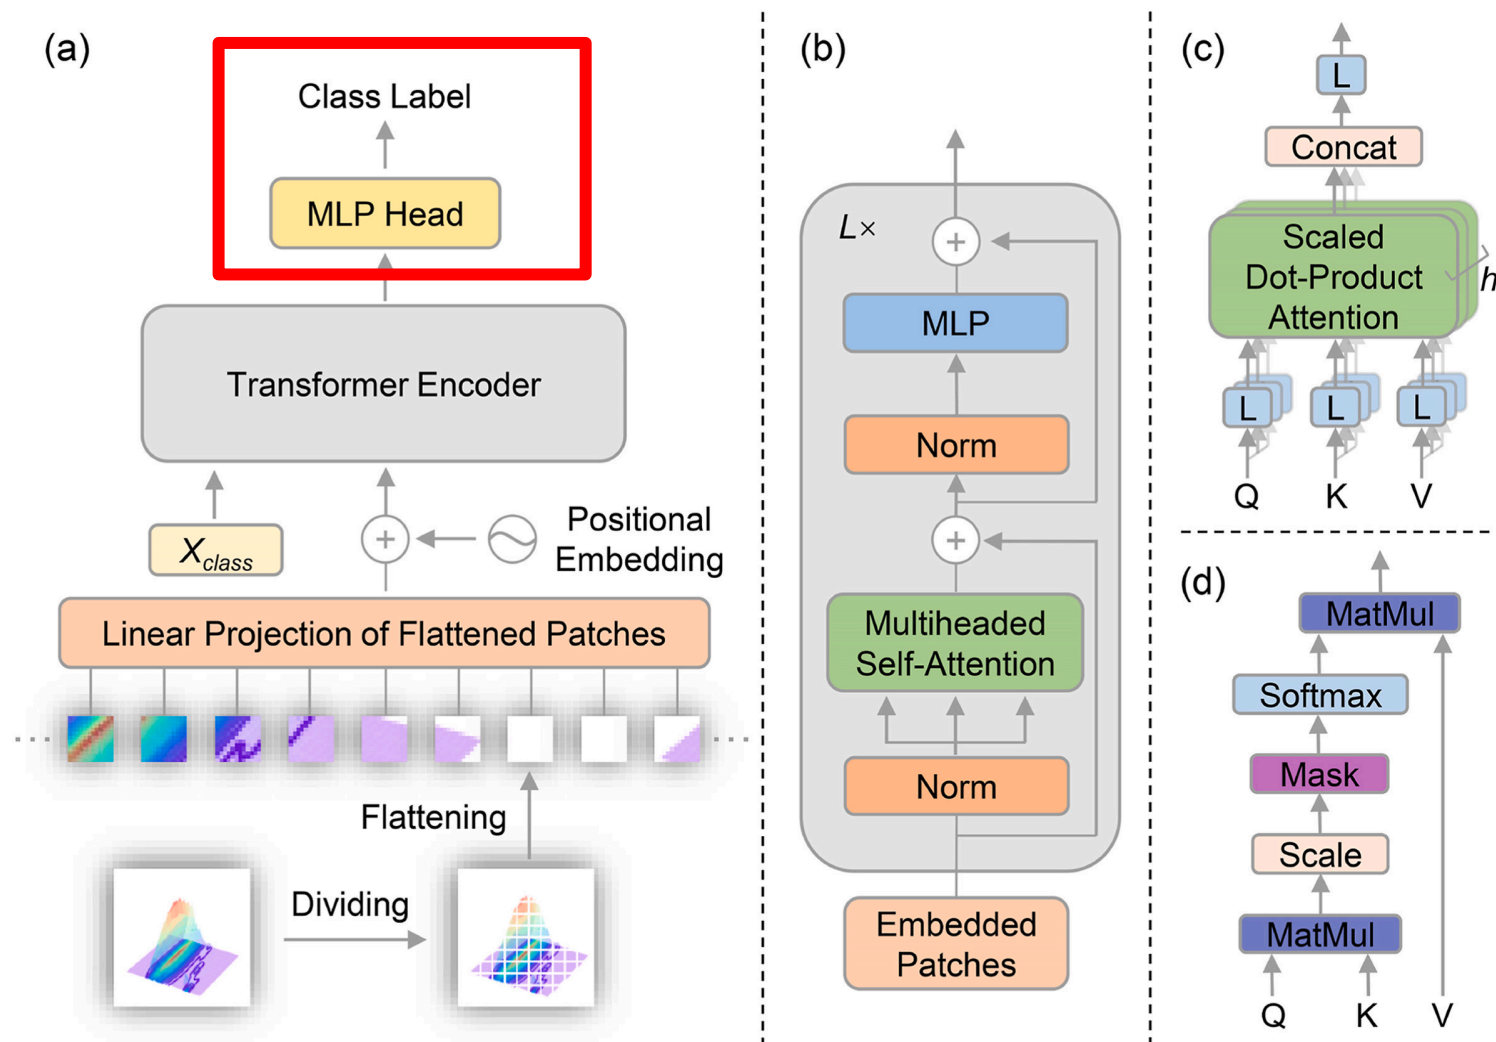


2.2 Multi-head Attention

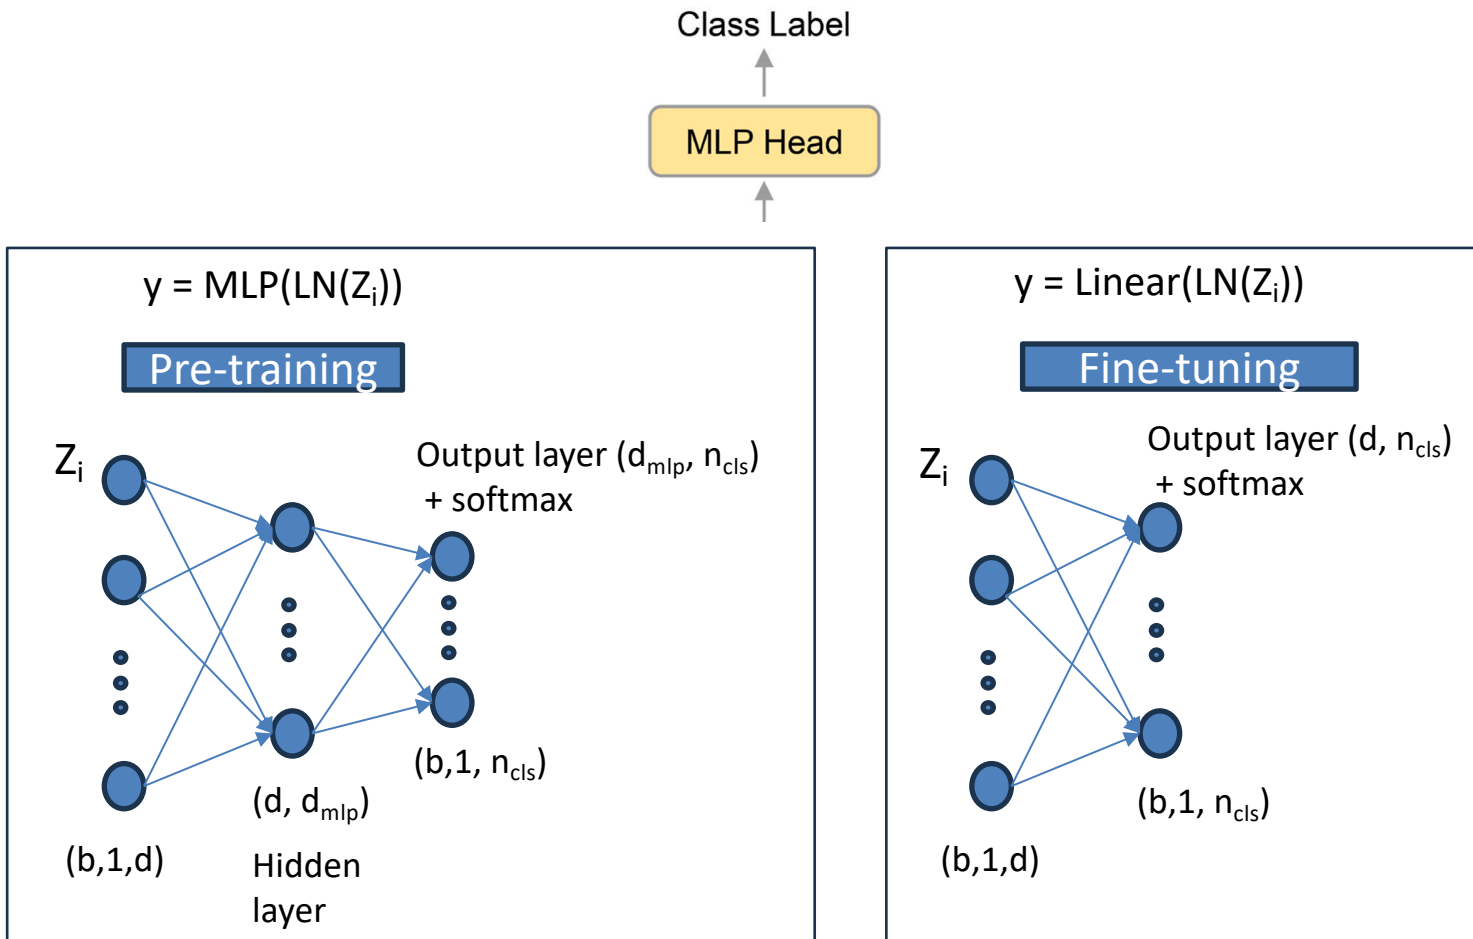
2. Transformer Encoding



3. Classification Head



3. Classification Head



- $[\text{cls}]$ token is used in the classification head.
- Pre-training - 2 layer of MLP used, hence 2 weight matrices
 - $W_h [d, d_{mlp}]$
 - $W_o [d_{mlp}, d]$
- Fine-tuning – single layer used, hence only 1 tensor $[d, n_{cls}]$
- Output: Probability associated with each of n_{cls} classes

Training Vision Transformer

Training Vision Transformers

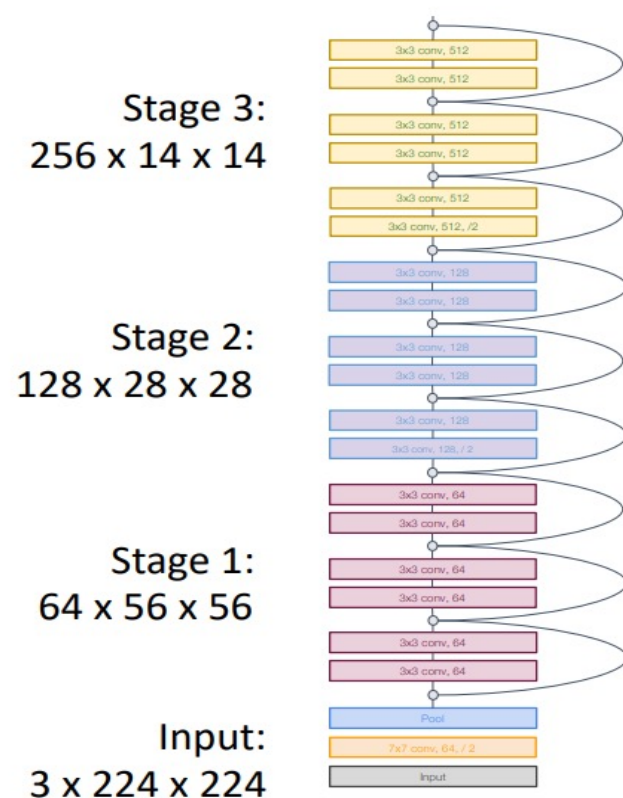
Model	Layers	Hidden size D	MLP size	Heads	Params
ViT-Base	12	768	3072	12	86M
ViT-Large	24	1024	4096	16	307M
ViT-Huge	32	1280	5120	16	632M

Applications of ViT

- Image Classification
- Image captioning
- Object Detection
- Semantics Segmentation
- Video Understanding
- Image Generation

Performance, Limitations and Future Improvements

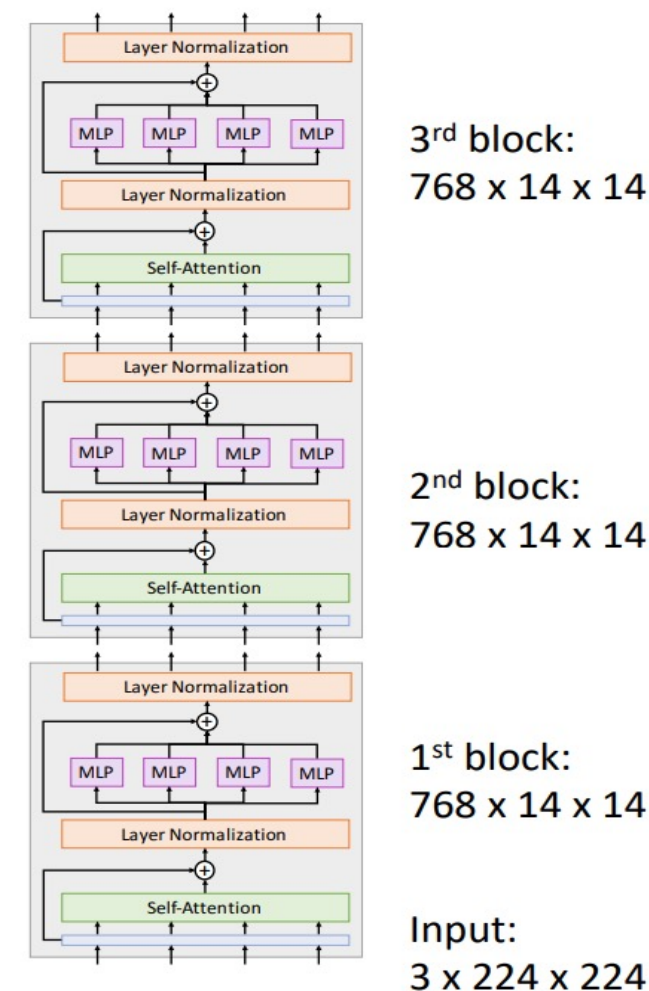
VIT vs CNN



In most CNNs (including ResNets), **decrease** resolution and **increase** channels as you go deeper in the network (Hierarchical architecture)

Useful since objects in images can occur at various scales

In a ViT, all blocks have same resolution and number of channels (Isotropic architecture)



Limitations of ViT

- Limited spatial information
- Computational Complexity
- Generalization to Diverse Data
- Memory and Computational Efficiency
- Continual Learning and Adaptability

Discussion and Conclusion

- **Superior Performance:**
 - Demonstrated exceptional performance, surpassing traditional CNNs in various image analysis tasks..
- **Global Context Modelling:**
 - capture global contextual information, enabling modeling of long-range dependencies and interactions between image elements.
- **Scalability and Adaptability:**
 - Scalability and Adaptability: ViTs are highly scalable and adaptable, making them suitable for diverse applications and datasets.

Future Improvements

- Enhanced attention mechanism for better long-range dependency modeling.
- Improved model efficiency with reduced computational complexity.
- Advancements in interpretability for better understanding of attention maps.
- Integration of multimodal information for enhanced performance.

References

- Han, Kai & Wang, Yunhe & Chen, Hanting & Chen, Xinghao & Guo, Jianyuan & Liu, Zhenhua & Tang, Yehui & Xiao, An & Xu, Chunjing & Xu, Yixing & Yang, Zhaohui & Zhang, Yiman & Tao, Dacheng. (2022). **A Survey on Vision Transformer.** IEEE Transactions on Pattern Analysis and Machine Intelligence. PP. 1-1. 10.1109/TPAMI.2022.3152247.
- C. Subakan, M. Ravanelli, S. Cornell, M. Bronzi and J. Zhong, "**Attention Is All You Need In Speech Separation,**" ICASSP 2021 - 2021 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), Toronto, ON, Canada, 2021, pp. 21-25, doi: 10.1109/ICASSP39728.2021.9413901.
- Dosovitskiy, Alexey & Beyer, Lucas & Kolesnikov, Alexander & Weissenborn, Dirk & Zhai, Xiaohua & Unterthiner, Thomas & Dehghani, Mostafa & Minderer, Matthias & Heigold, Georg & Gelly, Sylvain & Uszkoreit, Jakob & Houlsby, Neil. (2020). **An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale.**

Thankyou